

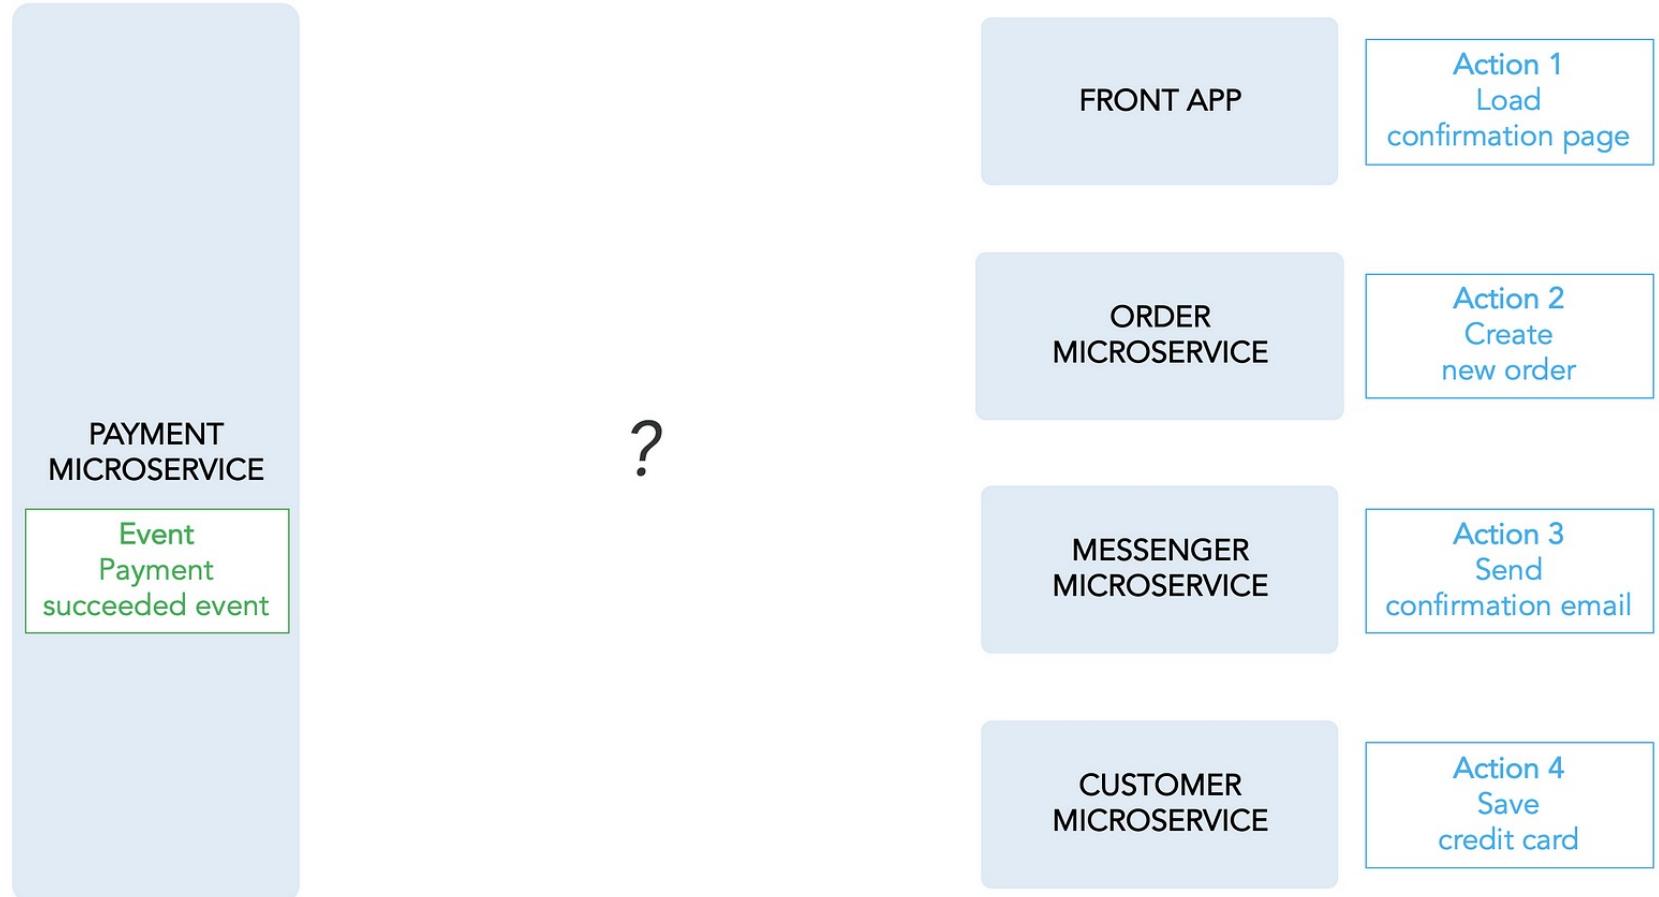
Event-driven systems

UA.DETI.IES

Scaling with Microservices

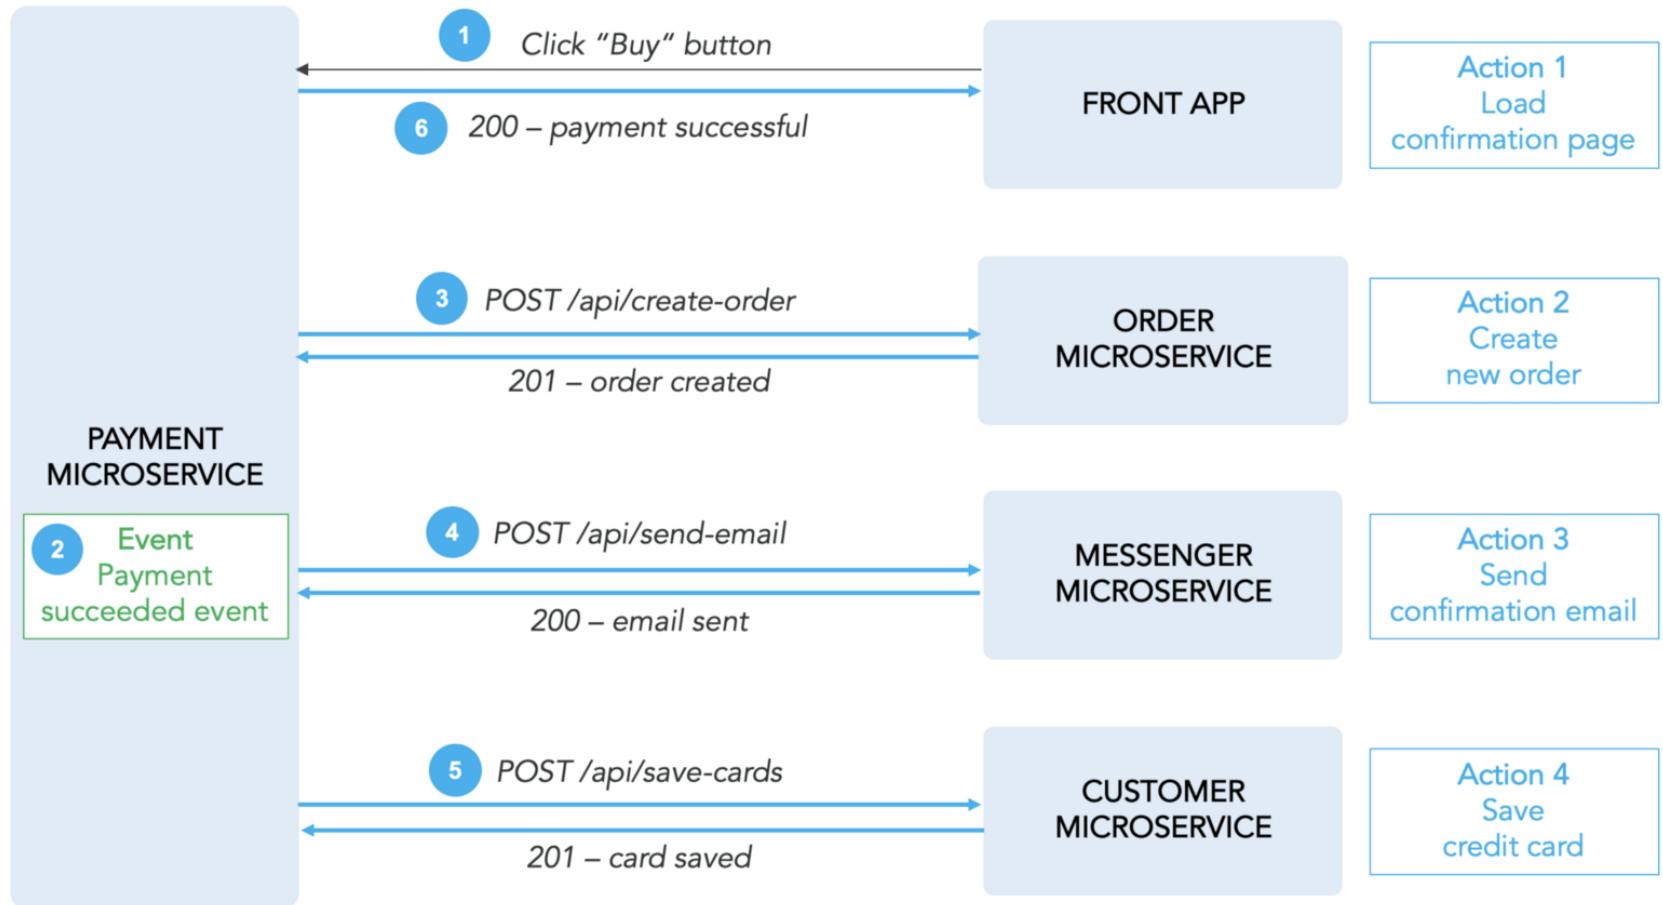
- ❖ **Application development** is currently based on concepts like **microservices** and **serverless** functions
 - We now live in a **cloud-native world**, and these models are a natural fit for distributed cloud-based environments
 - But simply building and deploying services and functions is not enough
- ❖ On its own, a single microservice does not accomplish much
 - **We also need a way to wire up those components**
 - i.e., to connect them so they can exchange data, forming a true application
 - In SE, this is one of the most important architectural decisions to make

How to notify every Microservice?

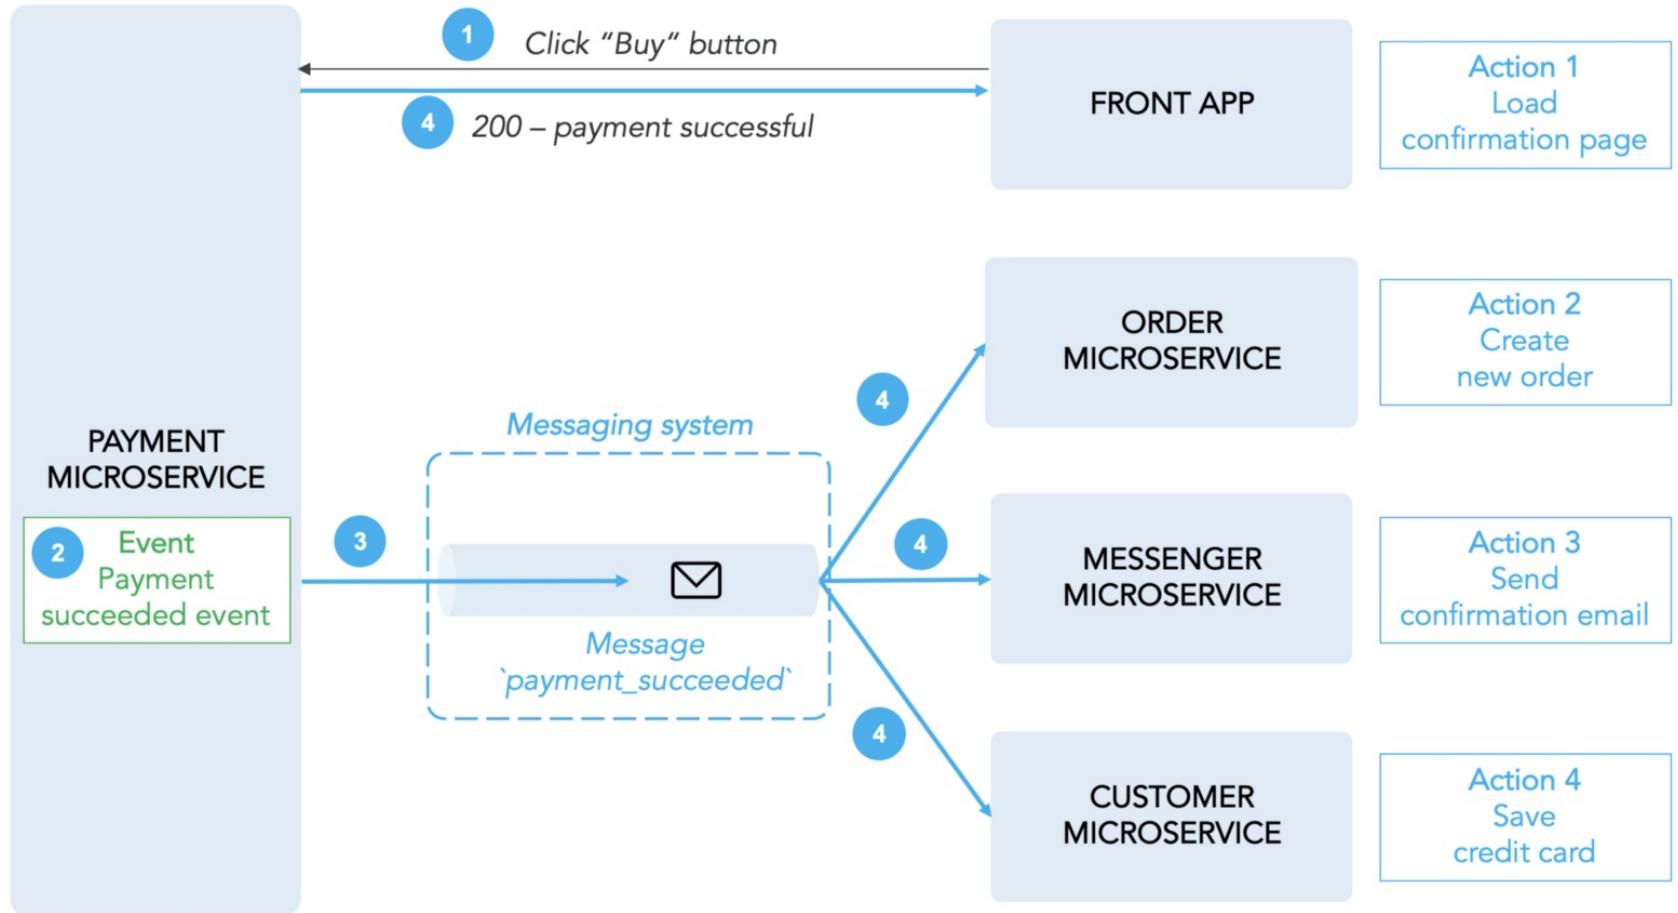


<https://blog.theodo.com/2019/08/event-driven-architectures-rabbitmq/>

Request-driven architecture



Event-driven architecture



Event-driven architecture

- ❖ For controlling these services, various mechanisms were developed over the years, such as message queues and **enterprise service buses** (ESBs).
 - E.g. **RabbitMQ**, **WSO2**.
- ❖ More recent offerings, the concept of streaming data have also emerged.
 - E.g. Apache **Kafka**
- ❖ This latter category is growing
 - Because streaming data is seen as a useful tool for implementing **event-driven architecture**
 - a software design pattern in which application data is modeled as streams of events, rather than as operations on static records.



What is an event?

- ❖ **Events** are things that happen, within a software system or, more broadly, during the operation of a business or other human process.
 - e.g., a sensor reports a temperature change, a user clicks their mouse, a customer deposits a check into a bank account.
- ❖ The concept of events in software systems closely aligns with how most of us think about our day-to-day lives.
 - Organizing around events makes it easier to develop business logic that accurately models real-world processes.
 - It helps reducing the number of one-to-one connections within a distributed system increasing the value of the microservices.
- ❖ An **event-driven architecture** allows generating, storing, accessing and reacting to these events.

Events ...



Events vs. queries and commands

Sou um consumidor! //

Quero ver coisas novas
não altero moda ...

❖ **Queries** are a request to look something up

- Unlike events or commands, queries are free of side effects; they leave the state of the system unchanged.

❖ **Commands** are actions → Producer

- Requests for some operation to be performed and will change the state of the system.
- Synchronous and typically indicate completion.

	Behavior/state change	Includes a response
Command	Requested to happen	Maybe
Event	Just happened	Never ↳ posso ignorar...
Query	None	Always

Event-driven patterns – notification

①

- é um evento!
- ❖ A service sends **events to notify** other systems of a change in its domain
 - For example, a user account service might send a notification event when a new login is created
 - ❖ What other systems choose to do with that information is largely up to them
 - The service that issued the notification just carries on with its business
 - ❖ Notification events usually do not carry much data
 - Resulting in a loosely coupled system with minimal network traffic spent on messaging

↓
Com data é o state transfer,

Event-driven patterns – state transfer

(2)

- ❖ A step up from simple notification, in this model the recipient of an **event also receives the data** it needs to perform further work
 - E.g., the user account service might issue an event that includes a data packet containing the new user's login ID, full name, hashed password, and other pertinent details.
↑ Produz alguma data
- ❖ This model can be appealing to developers familiar with **RESTful interfaces**.
 - But, depending on the complexity of the system, it can lead to a lot of data traffic on the network and data duplication in storage

Event-driven patterns – **sourcing**

③

Qual foi o serviço que originou? Não sei!

- ❖ The goal of **event-sourcing** is to represent every change of state in a system as an event, each recorded in chronological order *Guardar todos as alterações no sistema!*
!
- ❖ In so doing, the **event stream itself becomes the principal source of truth for the system**
 - E.g., it should be possible to “replay” a sequence of events to recreate the state of a SQL database at a given time
- ❖ This model presents a lot of possibilities, but it can be challenging to get right
 - particularly when events require participation from external systems

Event-driven advantages (over REST)

❖ Asynchronous

- Allows resources to move to the next task once a unit of work is complete
- Events are queued or buffered which prevents consumers from putting back pressure on producers or blocking them

❖ Loose Coupling

Desacoplado!

- Services operate independently, without knowledge of other services, including their implementation details and transport protocol
- Services under an event model can be updated, tested, and deployed independently and more easily

<https://dzone.com/articles/best-practices-for-event-driven-microservice-archi>

Event-driven advantages (over REST)

❖ Easy Scaling

- Since the services are decoupled and typically perform only one task, tracking bottlenecks and scaling a service is easier.

↳ Passo aumentar os consumidores facilmente

❖ Recovery Support

- Can recover lost work by “replaying” events from the past.

↑
Replay dos eventos,



<https://dzone.com/articles/best-practices-for-event-driven-microservice-archi>

Event-driven disadvantages

Overthinking

- ❖ They are easy to **over-engineer** by separating concerns that might be simpler when closely coupled
 - they can require significant upfront investment, and often result in additional complexity, service contracts or schemas, polyglot build systems, and dependency graphs
- ❖ **Complex data and transaction management**
 - Typically, do not support ACID transactions ❌
 - Systems must carefully handle inconsistent data between services, incompatible versions, duplicate events
- ❖ Even with these drawbacks, ...
 - An event-driven architecture is usually the better choice for enterprise-level microservice systems
 - The pros—scalable, loosely coupled, dev-ops friendly design—outweigh the cons

Event-driven Anti-Patterns

Não é sincrono ...

❖ Depending on Guaranteed Order and Delivery

- Events are asynchronous

- Including assumptions of order or duplicates will add complexity and will negate the key benefits of the event-based architecture.



❖ Premature Optimization

*{ Para sistemas pequenos
não compensa ... }*

- Most products start off small and grow over time

- Consider a simple architecture but include the necessary separation of concerns so that you can swap it out as your needs grow.

❖ Expecting Event-Driven to Fix Everything

- Event-driven architecture to fix all the problems

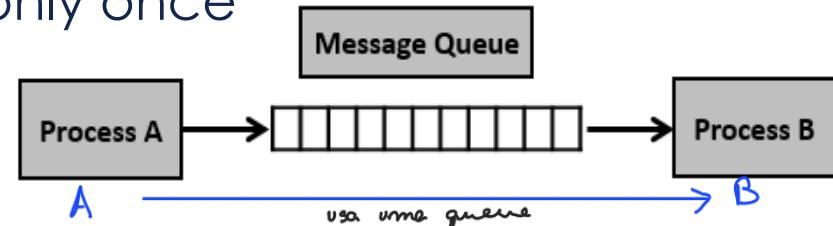
- It can't fix core problems such as a lack of automated testing, poor team communication, or outdated dev-ops practices.

Messaging systems

Messaging systems – models

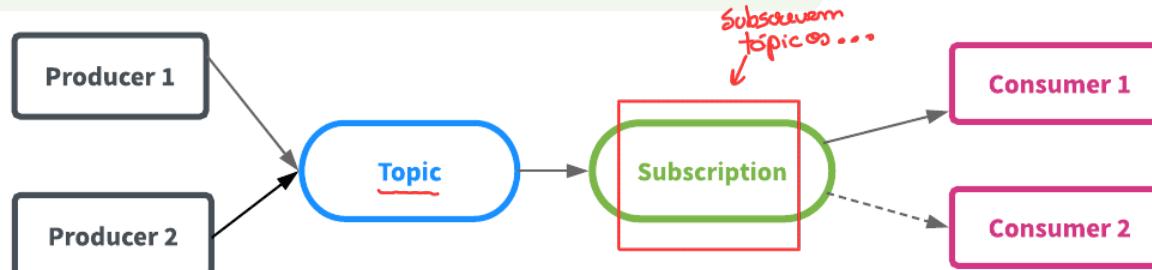
❖ Point to Point – Message queue

- Messages are sent to a queue to pre-defined receivers
- One-to-one relationship between sender and consumer
- Each message is consumed only once



❖ Publish-Subscribe

- Each message is published to a topic, and every application that subscribes to that topic gets a copy of all messages published to it.
↳ Normalmente com um retention time,
- Message producers are also known as publishers and consumers are known as subscribers.



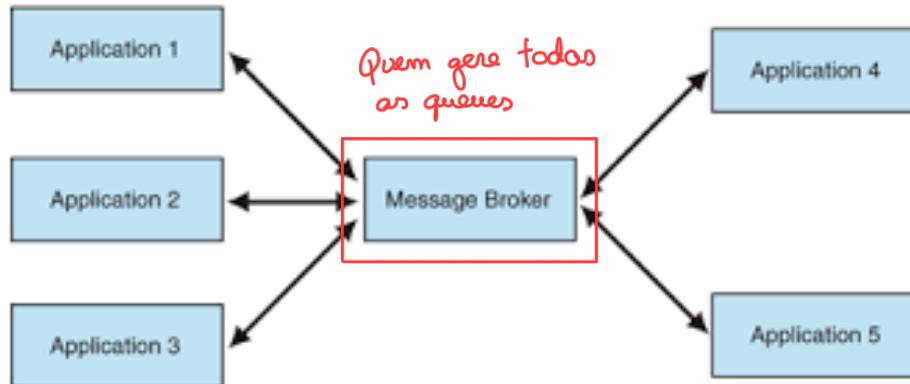
Messaging systems

❖ Managing the messages' flow

- Messages are “put into” a source queue
- They are then “taken from” a destination queue
- How/Who/What moves a message from a source queue to a destination queue?

❖ Queue Manager / **Message Broker**

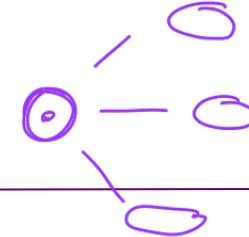
- Function as message-queuing “relay” that interact with distributed applications



Message broker

- ❖ Message Broker is built to extend MQ
 - it can understand the content of each message that it moves through the Broker.
- ❖ Message Broker can do the following:
 - divide the publisher and consumer
 - store and route the messages between services
 - Não esquecer* – converts between different transport protocols
 - identifies and distributes business events from disparate sources

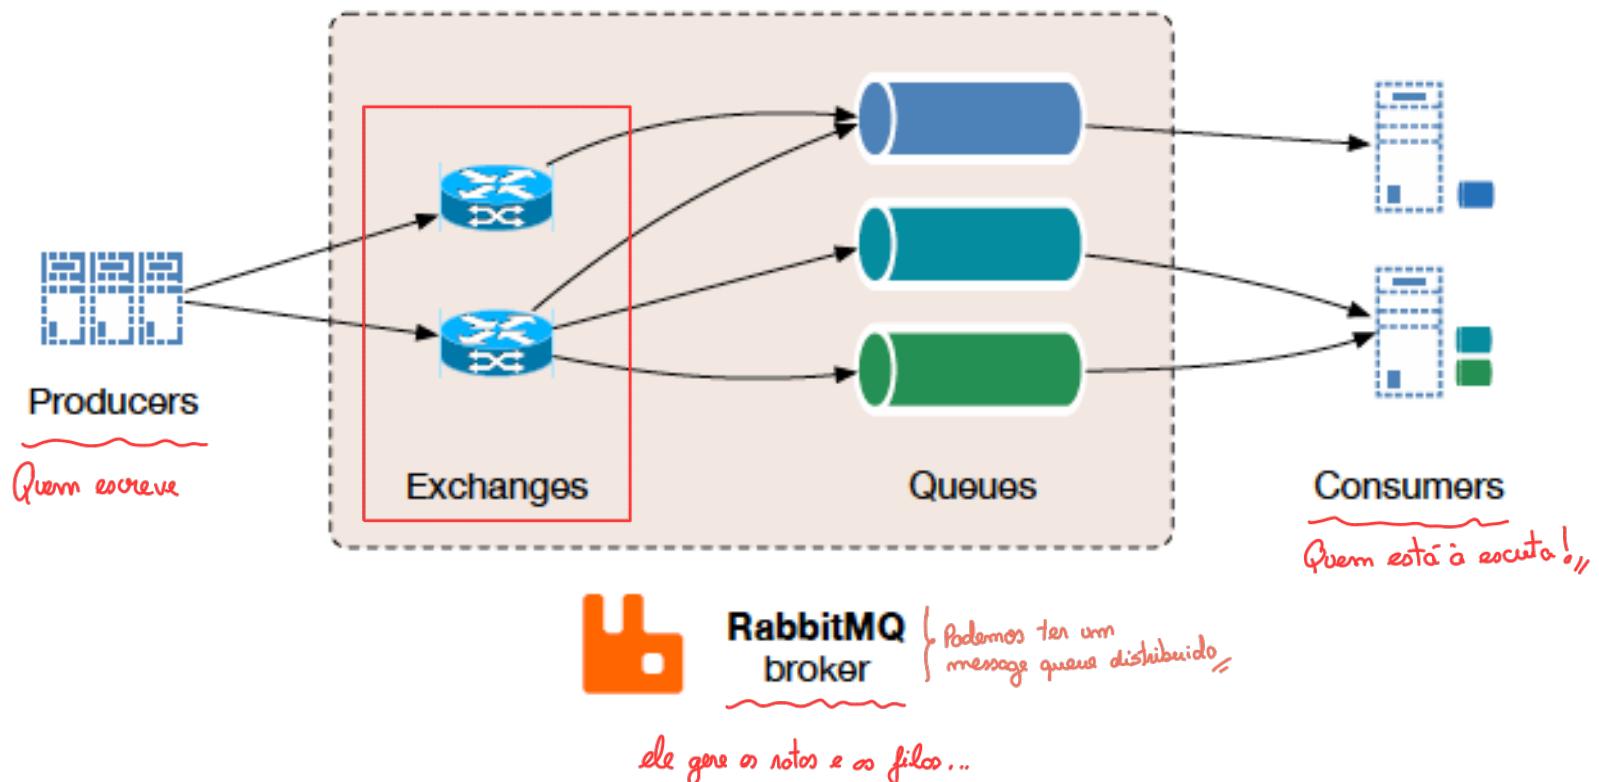
Message broker



- ❖ When is a message broker needed?
 - If we want to **control data feeds**, e.g., the number of registrations in a system
 - When the task is to **put data to several applications** and avoid direct usage of their API
 - When there is a need to **complete processes in a defined order** like a transactional system → *put in queue*
- ❖ There are many messaging tools...
 - E.g. **Apache ActiveMQ, RabbitMQ**
- ❖ ... and protocols
 - E.g. AMQP (Advanced Message Queuing Protocol), MQTT (MQ Telemetry Transport)

RabbitMQ

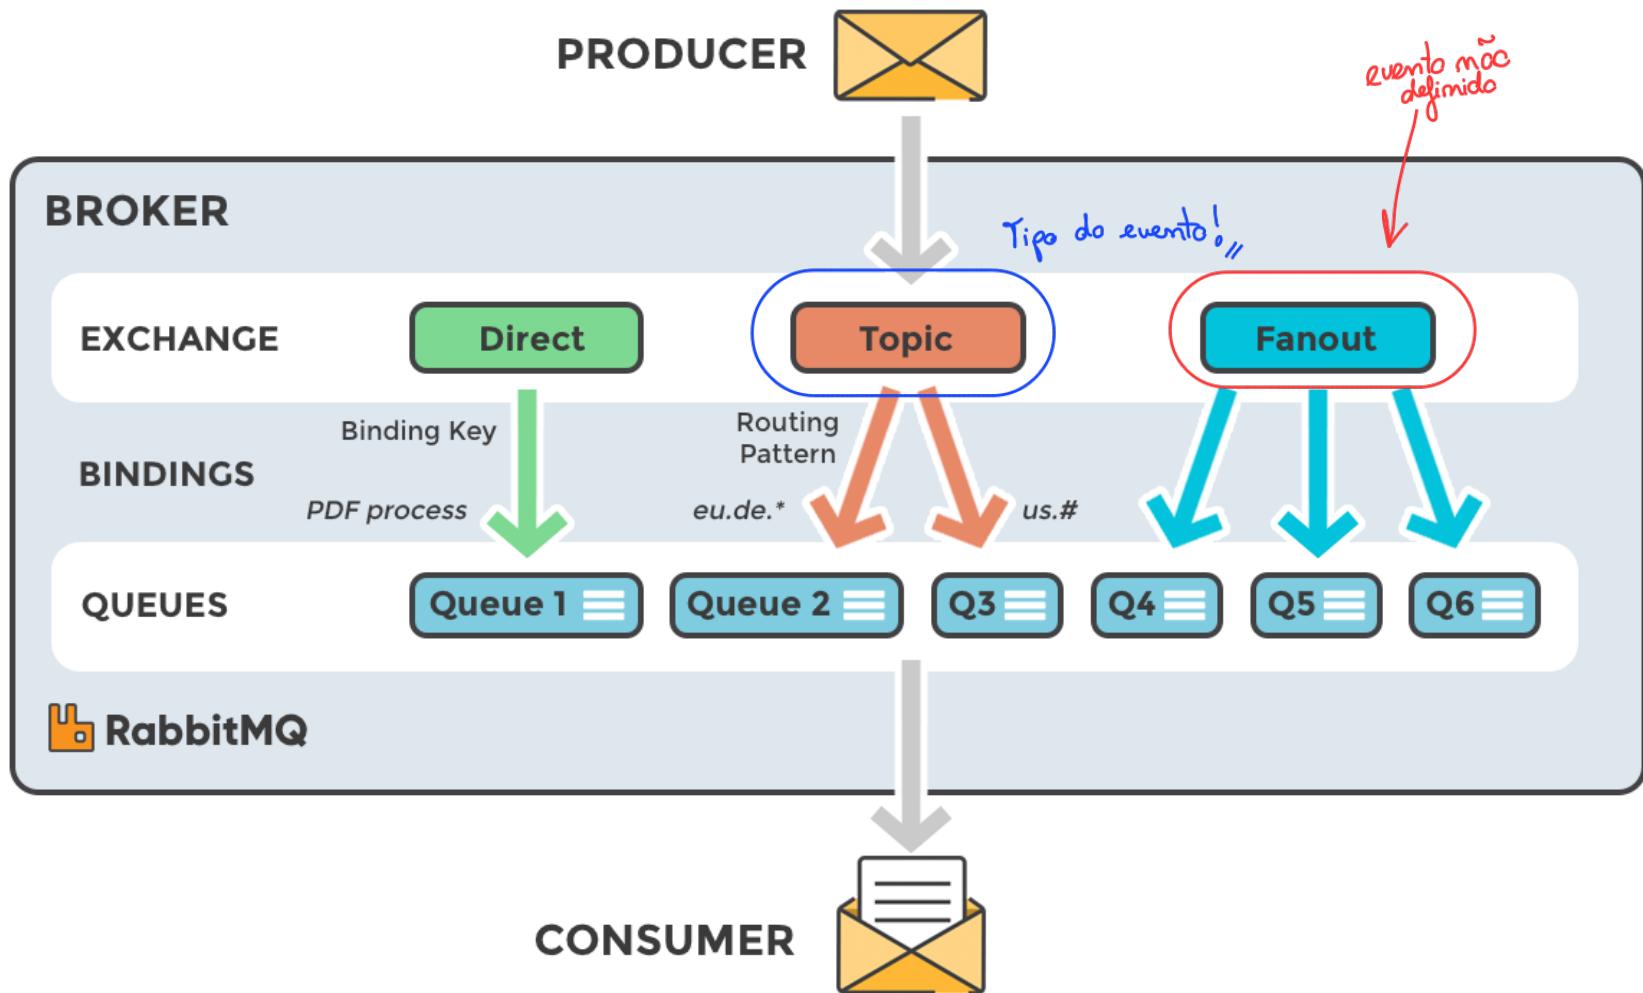
- ❖ RabbitMQ is one such open-source message broker software that implements AMQP.



RabbitMQ – Main concepts

- ❖ **Messages** contains attributes (like headers in a request) and a payload (the message content). *entidade no meio que faz o trânsito*
- ❖ Messages are published to an entity, **exchange**, which distribute the messages to **queues** (or Topics). *Quem recebe as mensagens*
- ❖ The rules for delivering the messages to the right queues are defined through
 - **bindings** (links between exchanges and queues), and
 - **routing keys** (a specific message attribute used for routing). *↳ quase como um tópico...*
- ❖ Messages stored in queues are
 - delivered continually to subscribers, or
 - fetched by consumers on demand.

RabbitMQ – Main concepts



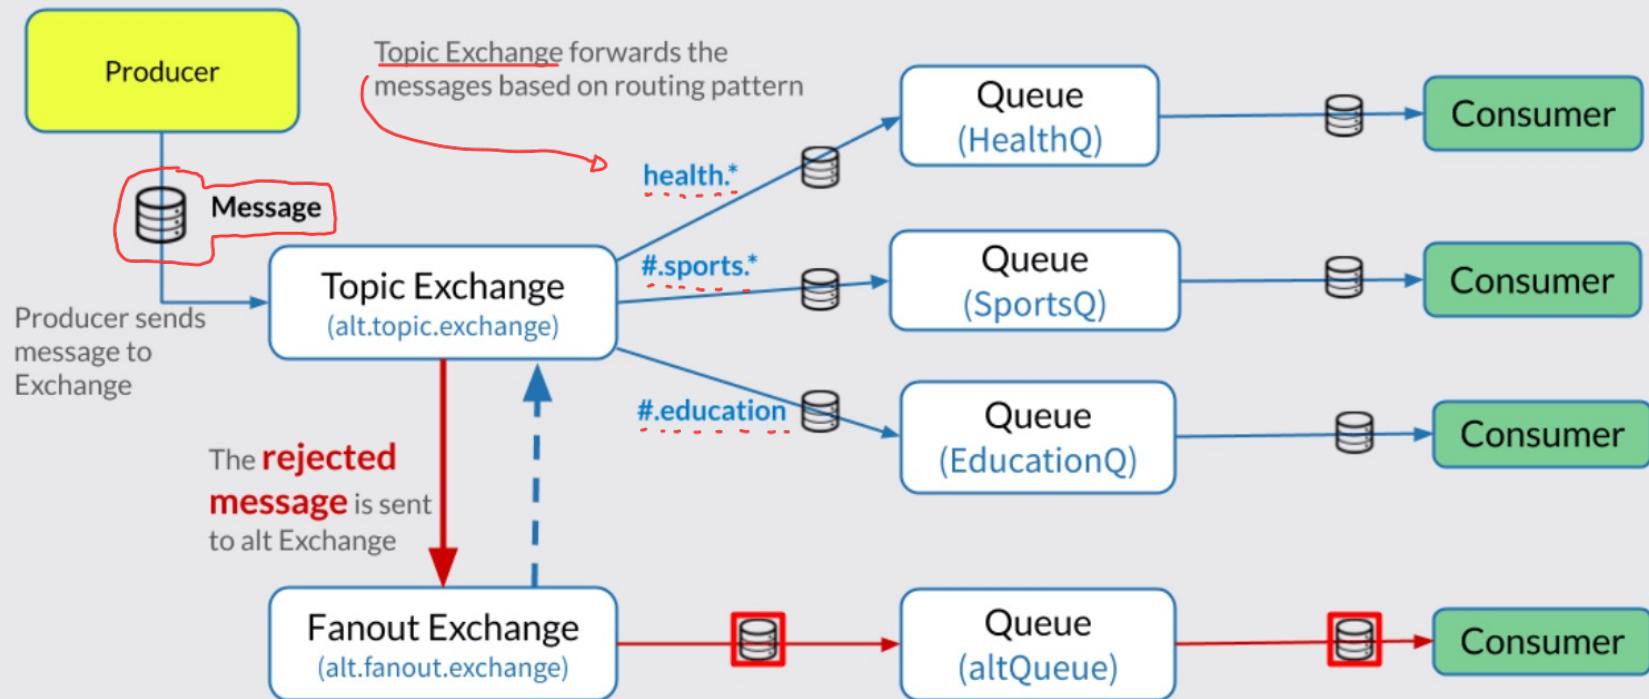
RabbitMQ – Exchange types (AMQP)

- ❖ **Direct Exchange** - It routes messages to a queue by matching **routing key** equal to **binding key**.
- ❖ **Topic Exchange** – It routes messages to multiple queues by a partial matching of a routing key. It uses patterns to match the routing and binding key.
- ❖ **Fanout Exchange** - It ignores the routing key and sends message to all the available queues. *→ Manda para todos!..*
- ❖ **Headers Exchange** – It uses message header instead of routing key.
- ❖ **Default (Nameless) Exchange** - It routes the message to queue name that exactly matches with the routing key.

Alertas! por
exemplo...

RabbitMQ – Exchange types

Alternate Exchange in RabbitMQ



Spring Boot – Messaging with RabbitMQ

- ❖ Set up the RabbitMQ Broker
- ❖ Spring Initializr
- ❖ E.g. RabbitMQ dependency

```
<!-- https://mvnrepository.com/artifact/org.springframework.boot/spring-boot-starter-amqp -->
<dependency>
<groupId>org.springframework.boot</groupId>
<artifactId>spring-boot-starter-amqp</artifactId>
<version>2.3.2.RELEASE</version>
</dependency>
```

- ❖ <https://spring.io/guides/gs/messaging-rabbitmq/>

Example excerpt ...

```
@SpringBootApplication
public class MessagingRabbitmqApplication {
    static final String topicExchangeName = "spring-boot-exchange";
    static final String queueName = "spring-boot";

    @Bean
    Queue queue() {
        return new Queue(queueName, false);
    }

    @Bean
    TopicExchange exchange() {
        return new TopicExchange(topicExchangeName);
    }

    @Bean
    Binding binding(Queue queue, TopicExchange exchange) {
        return BindingBuilder.bind(queue).to(exchange).with("foo.bar.#");
    }
}
```

Hello world – Sender

```
public class Sender {  
  
    @Autowired  
    private RabbitTemplate template;  
  
    @Autowired  
    private Queue queue;  
  
    @Scheduled(fixedDelay = 1000, initialDelay = 500)  
    public void send() {  
        String message = "Hello World!";  
        this.template.convertAndSend(queue.getName(), message);  
        System.out.println(" [x] Sent '" + message + "'");  
    }  
}
```

Hello world – Receiver

```
@RabbitListener(queues = "hello")
public class Receiver {

    @RabbitHandler
    public void receive(String in) {
        System.out.println(" [x] Received '" + in + "'");
    }
}
```

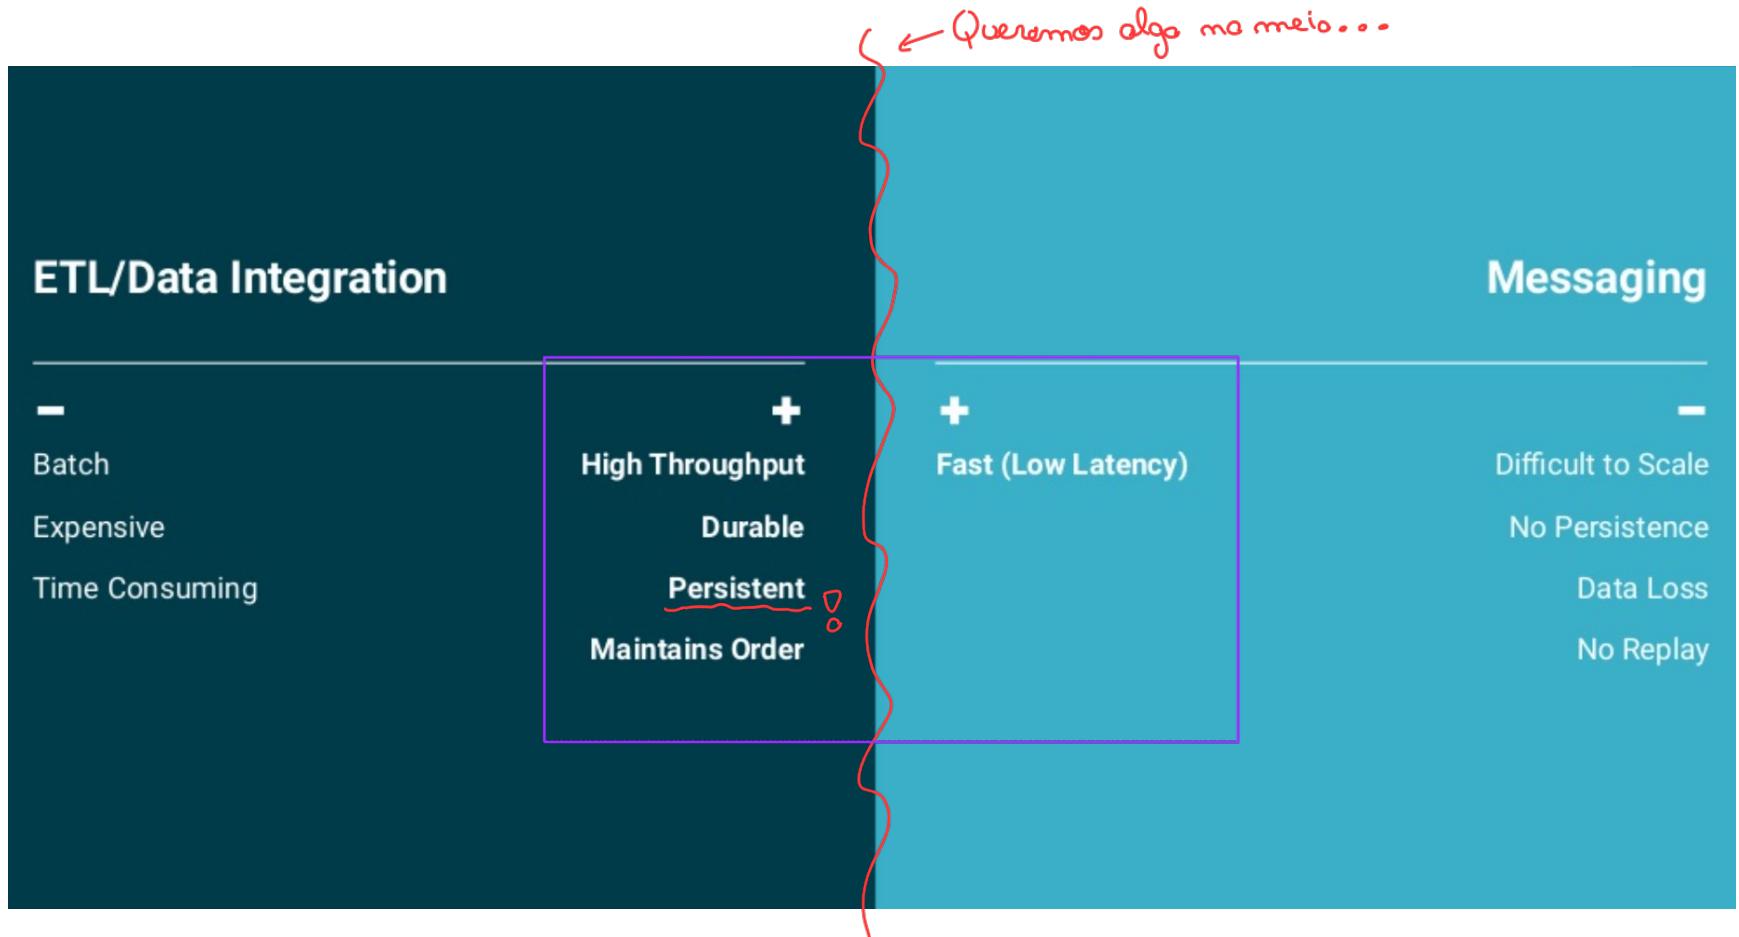
❖ Tutorial:

- <https://www.rabbitmq.com/tutorials/tutorial-one-spring-amqp.html>

Event streaming systems

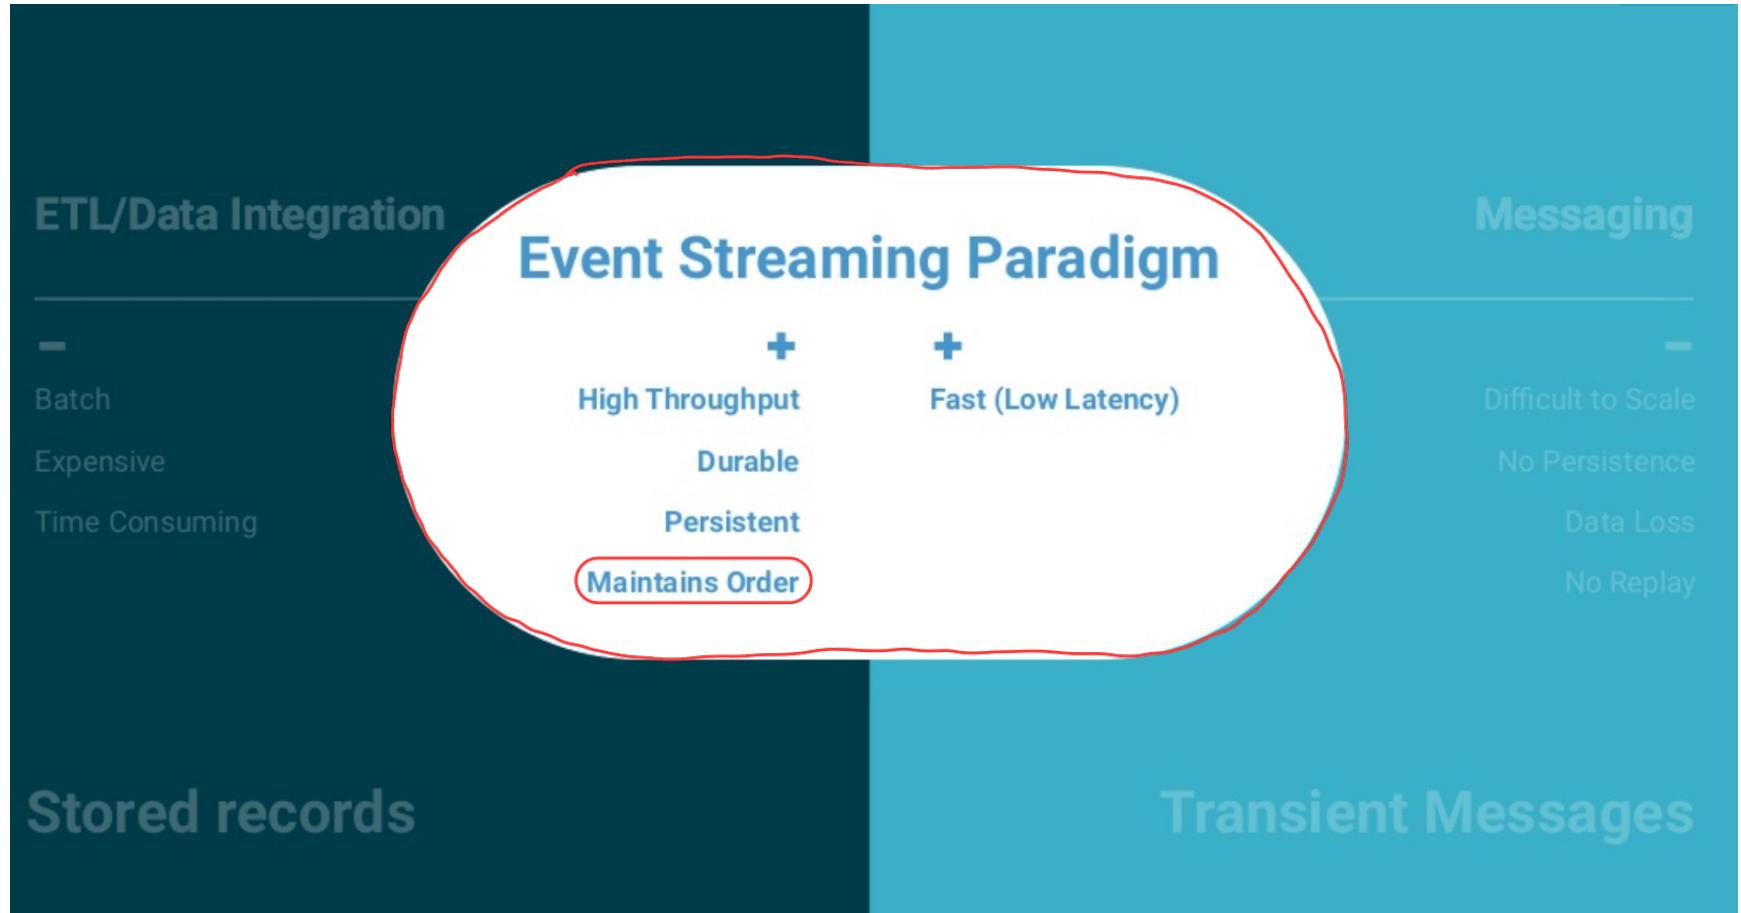
Apache Kafka

Stored records vs Messaging



<https://pt.slideshare.net/ConfluentInc/what-is-apache-kafka-and-what-is-an-event-streaming-platform>

The Event Streaming Paradigm



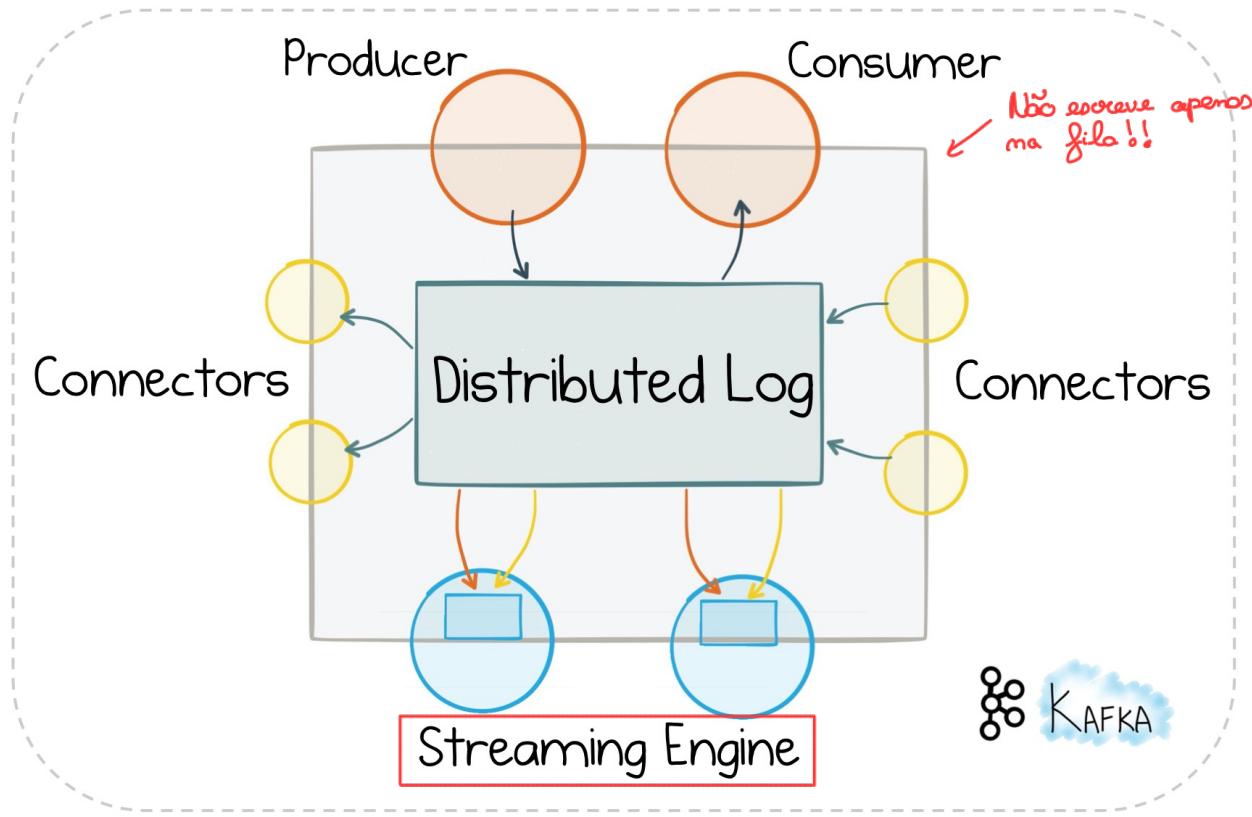
The Event Streaming Paradigm

To rethink data as not stored records or transient messages, but instead as a continually updating stream of events

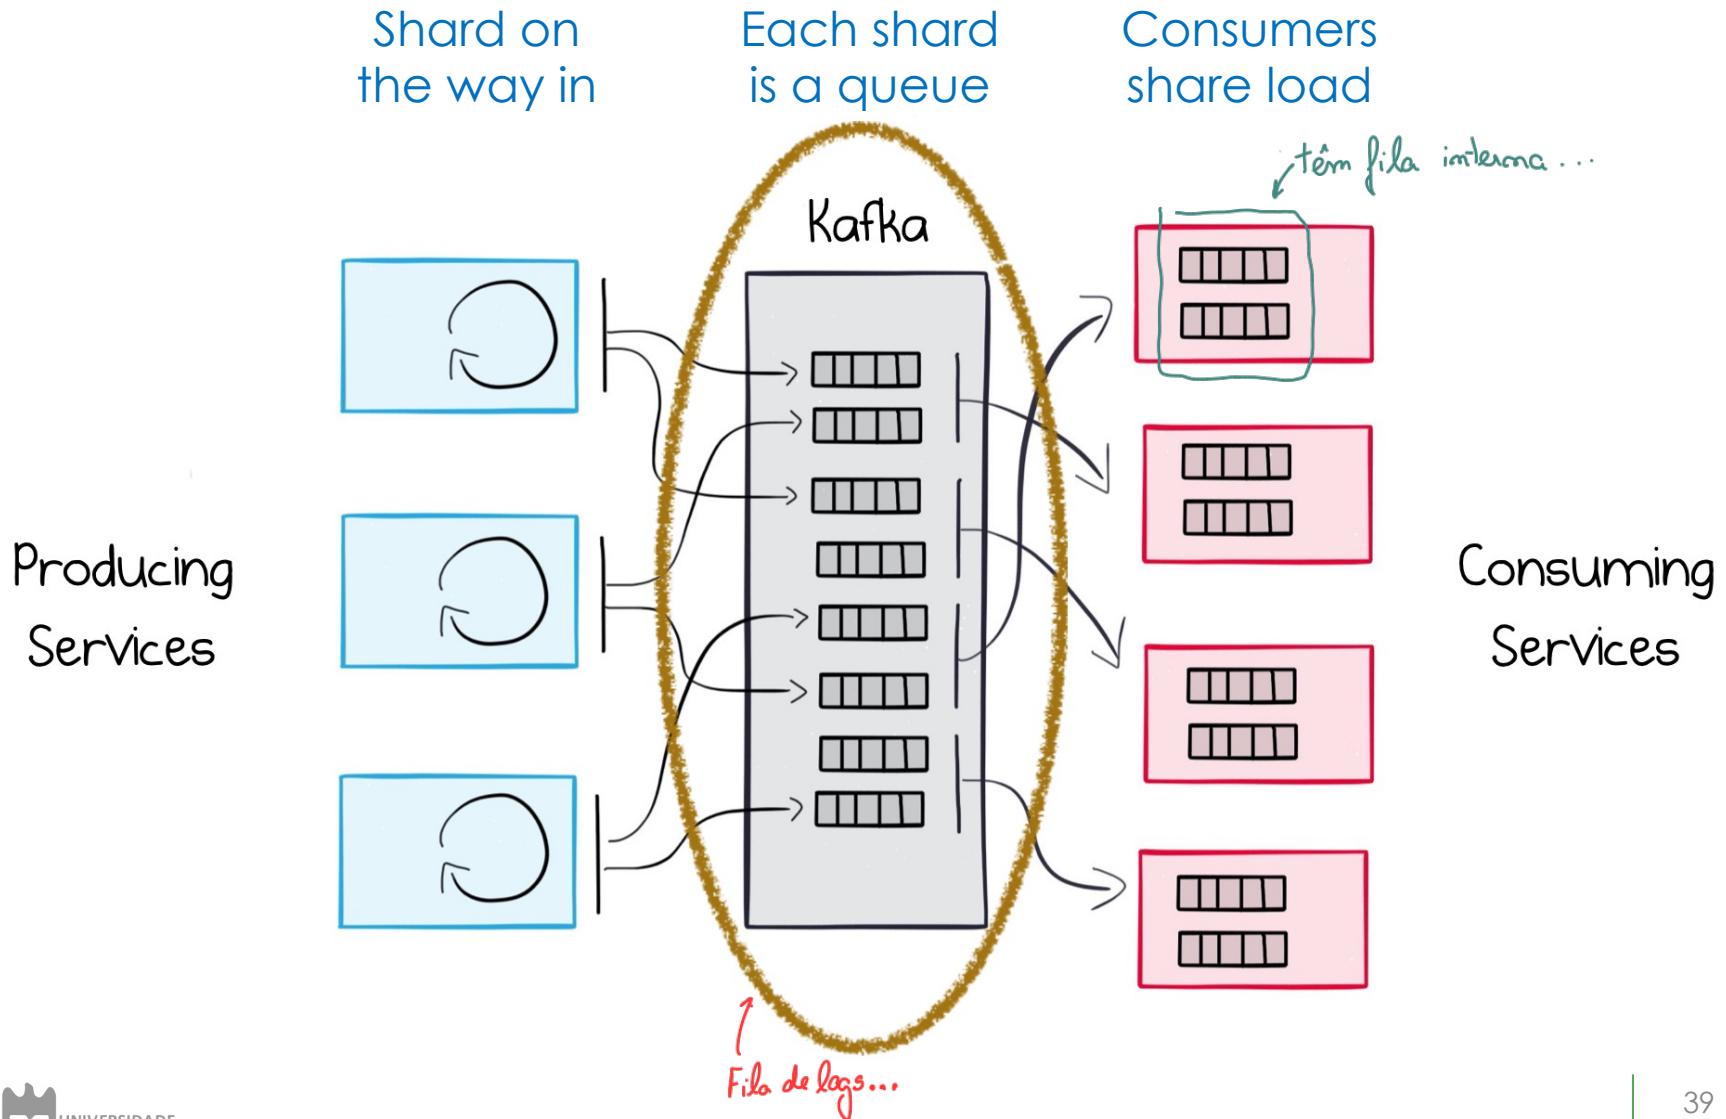
<https://pt.slideshare.net/ConfluentInc/what-is-apache-kafka-and-what-is-an-event-streaming-platform>

Apache Kafka

- ❖ Apache Kafka is made of distributed, immutable, append-only commit logs



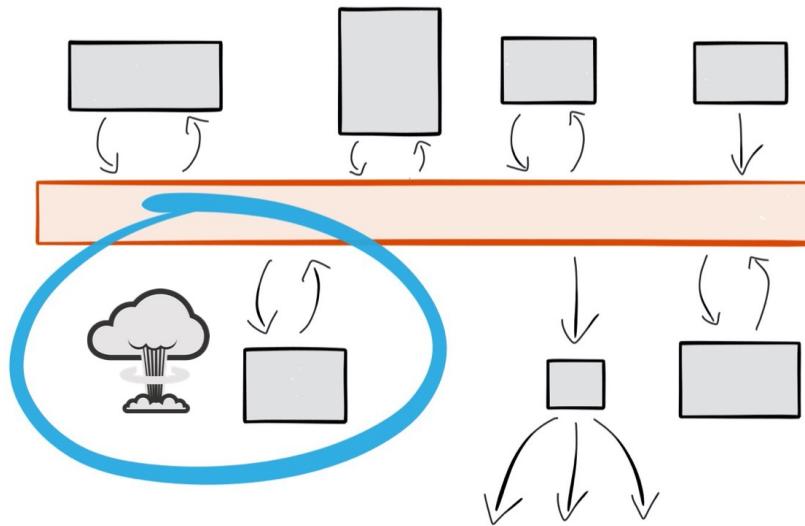
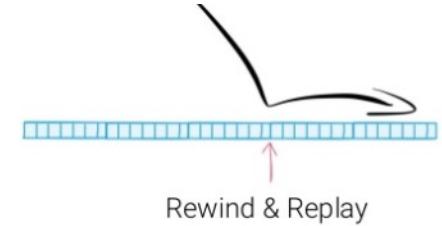
The Distributed Log



The Distributed Log

❖ The Service Backbone

- Scalable/Load Balanced, Fault Tolerant, Concurrent, Strongly Ordered, Stateful



❖ A place to keep data-on-the-outside

- When sending data across services, it is outside the normal "trust" boundary.

Events and topics

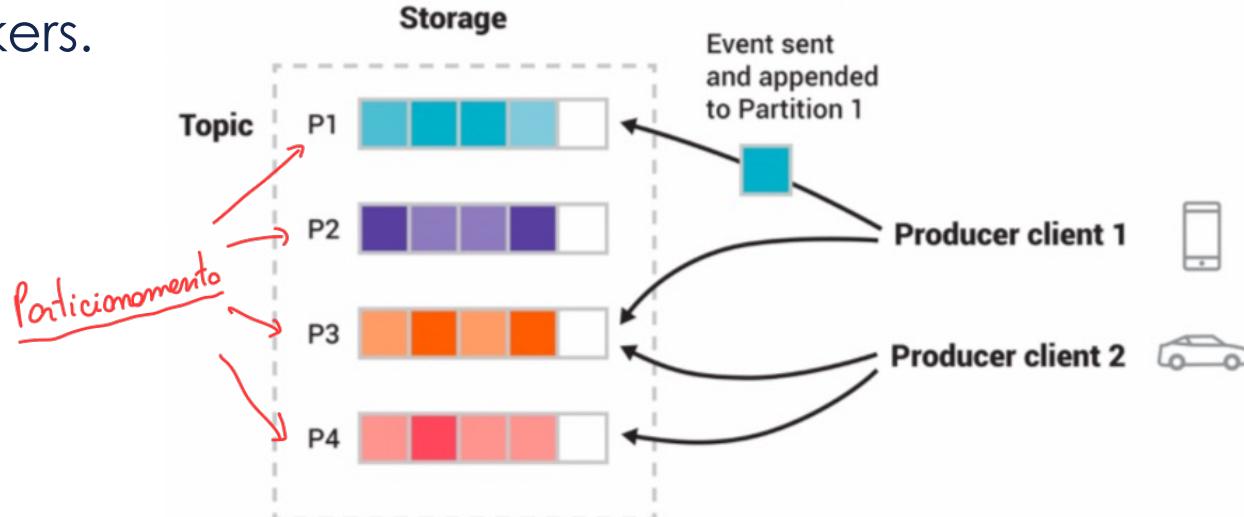
- ❖ Each **event** has a key, value, timestamp, and optional metadata headers.

Event key: "Alice"

Event value: "Made a payment of \$200 to Bob"

Event timestamp: "Jun. 25, 2020 at 2:06 p.m."

- ❖ Events are organized and durably stored in **topics**
- ❖ Topics are **partitioned**
 - a topic is spread over "buckets" located on different Kafka brokers.



Kafka APIs

- ❖ **Admin API** to manage and inspect topics, brokers, and other Kafka objects.
- ❖ **Producer API** to publish (write) a stream of events to one or more Kafka topics.
`KafkaProducer<K, V> (implements Producer<K, V>)`
- ❖ **Consumer API** to subscribe to (read) one or more topics and to process the stream of events produced to them.
`KafkaConsumer<K, V> (implements Consumer<K, V>)`
- ❖ The **Kafka Streams API** to implement stream processing applications and microservices.
`KStream<K, V>`
 - To process event streams, including transformations, stateful operations like aggregations and joins, windowing, processing based on event-time, and more.
- ❖ The **Kafka Connect API** to build and run import/export connectors to external systems and applications.
 - Kafka community already provides hundreds of ready-to-use connectors.

Event-driven patterns

❖ Previously...

- Event notification, Event-carried state transfer, Event-sourcing

Apenas a notificação
↓

com valores para
dizer ao destinatário
↓

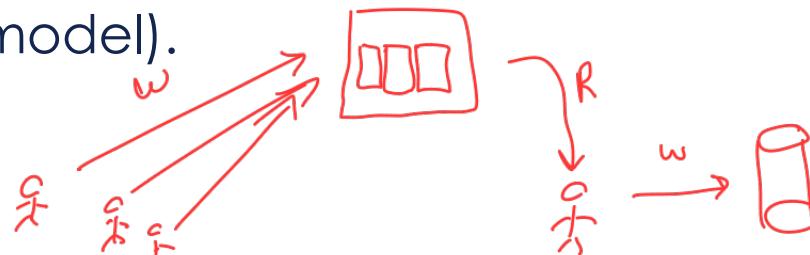
As mudanças do
estado são todos
guardados na
queue

❖ In **Event Sourcing**, events are a core element – the source of truth.

- Being stored, immutably, in the order they were created in, the event log expresses exactly what the system did.

❖ **Command Query Response Segregation (CQRS)** is a natural progression from this.

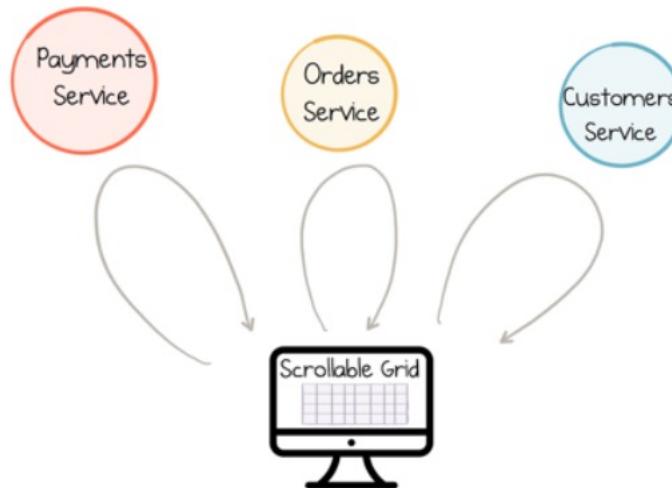
- Decoupling writing and reading operations.
- As a simple example, we might write events to Kafka (write model), read them back, and then push them into a database (read model).



Example

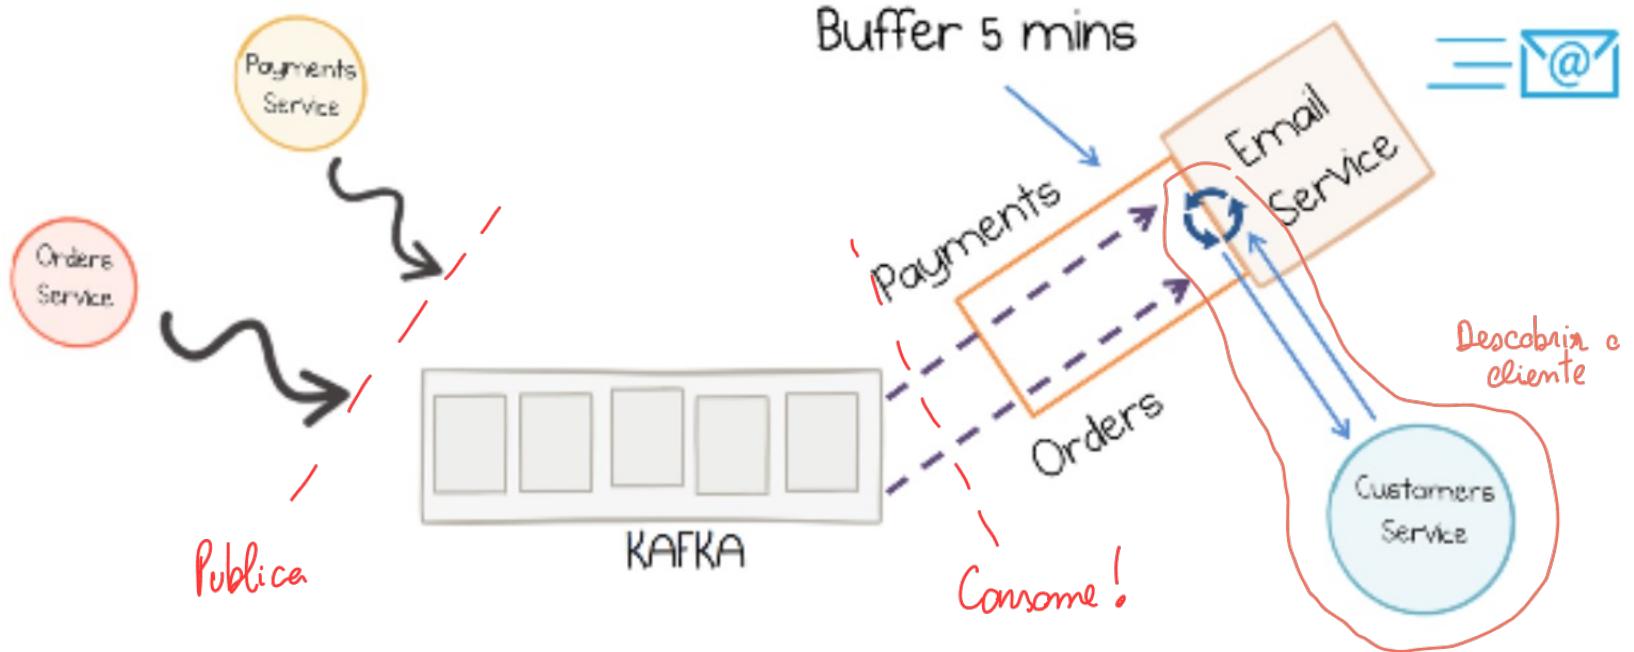
- ❖ Say we have an email service that listens to an event stream of orders and then sends confirmation emails to users once they complete a purchase.
 - This requires information about both the order as well as the associated payment. Such an email service might be created in several different ways

The Request Response Way



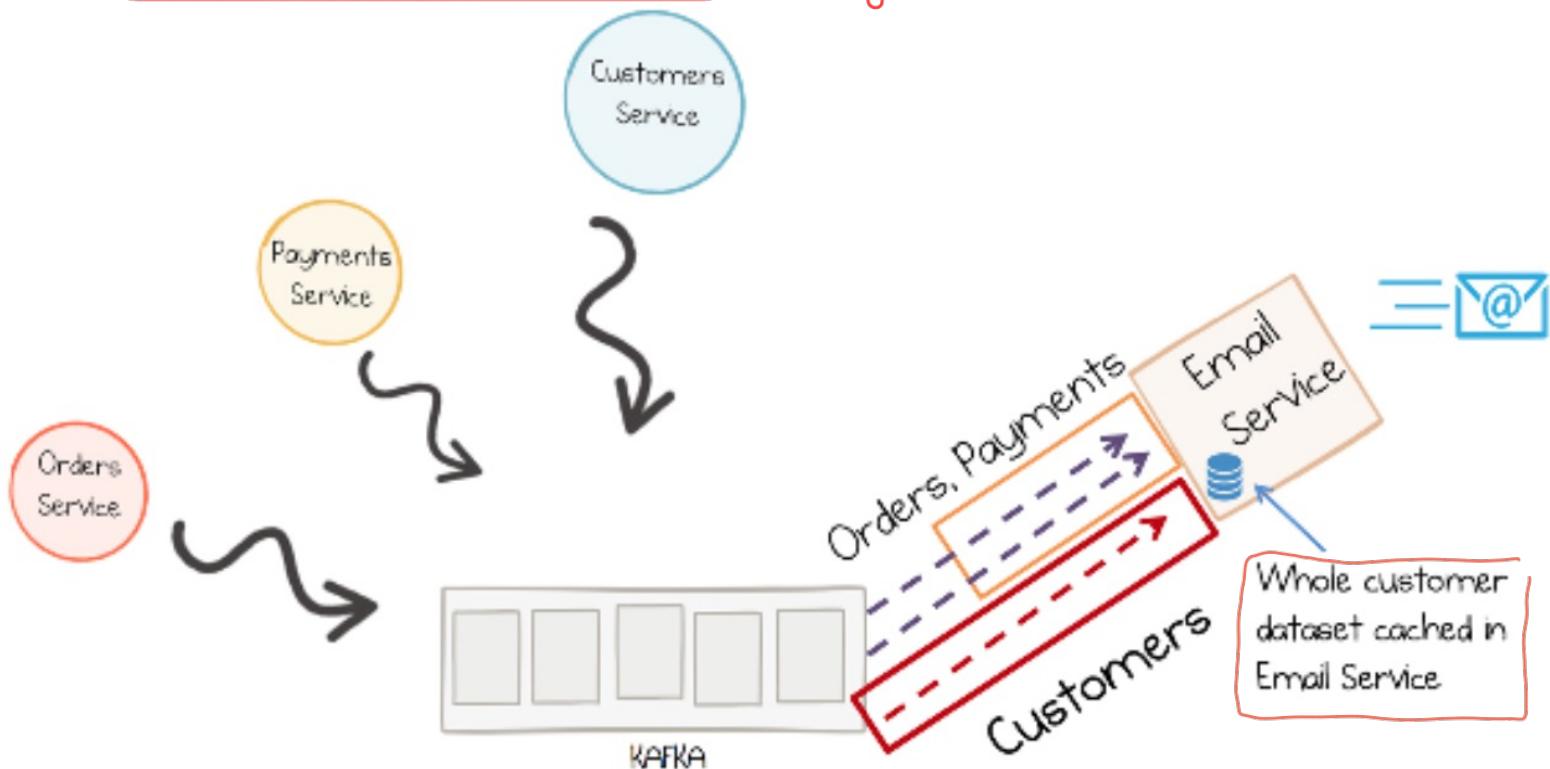
Stateless Streaming approach

- ❖ Example: A stateless streaming service that looks up reference data in another service at runtime



Stateful Streaming approach

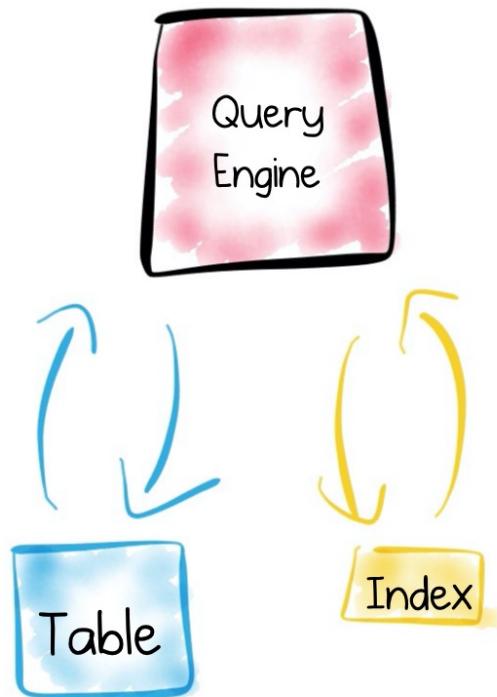
- ❖ Example: A **stateful streaming** service that replicates the Customers topic into a local table, held inside the Kafka Streams API. *← stateful*



Stateful Streaming approach

- ❖ Being stateful comes with some challenges:
 - when a new node starts, it must load all stateful components (i.e., state stores)
- ❖ **Kafka Streams**, for instance, provides three mechanisms to simplify stateful:
 - It uses a technique called **standby replicas**, which ensure that for every table or state store on one node, there is a replica kept up to date on another.
 - **Disk checkpoints** are created periodically so that, should a node fail and restart, it can load its previous checkpoint.
 - Finally, **compacted topics** are used to keep the dataset as small as possible. This acts to reduce the load time for a complete rebuild should one be necessary.

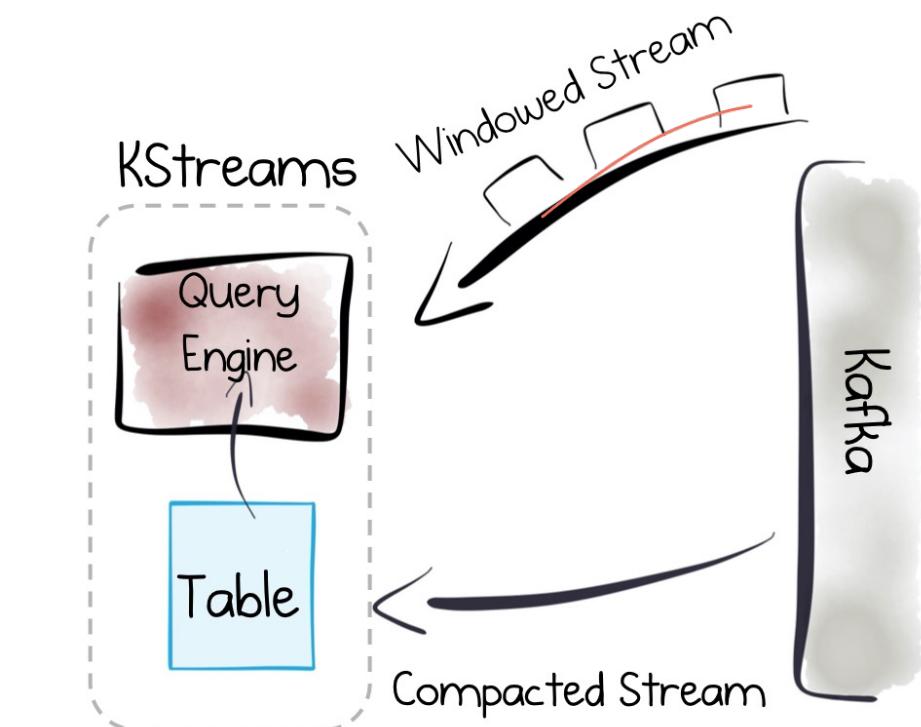
A Stateful Stream Processing



Database

Finite source

VS



Stateful Stream Processor

Infinite & Finite source

Messaging in Spring

- ❖ **Spring Cloud Stream** helps fully abstract code from the underlying messaging engine
- ❖ It supports a variety of binder implementations such as:
 - RabbitMQ
 - Apache Kafka
 - Kafka Streams
 - Amazon Kinesis
 - Google PubSub (partner maintained)
 - Solace PubSub+ (partner maintained)
 - Azure Event Hubs (partner maintained)
 - Apache RocketMQ (partner maintained)

Messaging in Spring

❖ Spring AMQP

- It applies core Spring concepts to the development of AMQP-based messaging solutions.
- The project consists of two parts; `spring-amqp` is the base abstraction, and `spring-rabbit` is the RabbitMQ implementation.

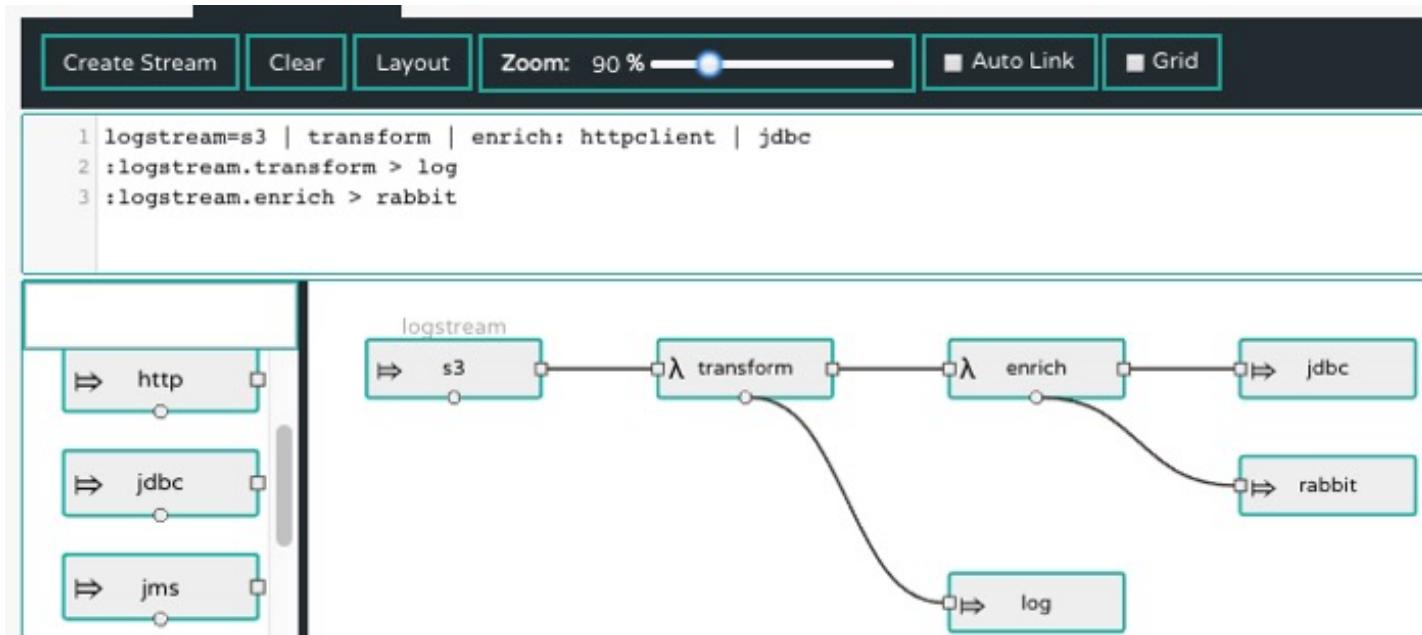
❖ Spring for Apache Kafka

- Simple snippet

```
@Configuration(proxyBeanMethods = false)
@EnableKafkaStreams
public static class KafkaStreamsExampleConfiguration {
    @Bean
    public KStream<Integer, String> kStream(StreamsBuilder streamsBuilder) {
        KStream<Integer, String> stream = streamsBuilder.stream("ks1In");
        stream.map((k, v) -> new KeyValue<>(k, v.toUpperCase())).to("ks10ut",
            Produced.with(Serdes.Integer(), new JsonSerde<>()));
        return stream;
    }
}
```

Messaging in Spring

- ❖ **Spring Cloud Data Flow** allows to create and orchestrate data pipelines, e.g., data ingest, real-time analytics, and data import/export.



Summary

- ❖ Event-driven architecture is gaining in popularity, and with good reason.
 - From a technical perspective, it provides an effective method of wiring up microservices.
 - The interest in serverless functions - such as AWS Lambda, Azure Functions, or Knative - is growing, and these are inherently event-based.
 - Moreover, when coupled with modern streaming data tools like Apache Kafka, event-driven architectures become more versatile, resilient, and reliable than with earlier messaging methods.
- ❖ But perhaps the most important “feature” of the event-driven pattern is that it models how businesses operate in the real world.

Resources & Credits



- ❖ Designing Event-Driven Systems
Ben Stopford, O'Reilly

- ❖ The Optimal RabbitMQ Guide, Lovisa Johansson, CloudAMQP

- ❖ Spring messaging projects
 - <https://spring.io/projects/>

- ❖ Kafka, RabbitMQ, etc..
 - <https://kafka.apache.org>
 - <https://www.rabbitmq.com>